

Parallelized PDE Solver Using Adaptive Meshing and User Supplied Domain and Boundary Conditions

Laurel Koenig, Isaac Shirk

Abstract

Numerical methods for computing solutions to partial differential equations (PDEs) have been well developed over the years due to their being ubiquitous in nearly all fields of science. No matter the field, be it manufacturing, to rocket science, PDEs need to be computed accurately and efficiently. However, numerically solving PDEs in two and three dimensions typically incurs a great computational cost, leading to impractically long computation times. This project uses a quadtree to create a variable grid and then balance the workload to parallelize the computation process. This results in some speedup, but not to the expected degree.

I. INTRODUCTION

THE Focus of this project is to create a tool that can produce solutions to PDEs on domains with boundary conditions provided by users. The domains include rectangular and circular two dimensional areas, with edge conditions being set to constant values, or a constant derivative. The user defines the shape of the domain, the boundary conditions, and the initial state, and the solver will then compute the solution for some time in the future.

Often, when numerically solving PDEs for physical applications, a very finely discretized spatial grid is necessary to obtain small errors. However, increased grid resolution comes at the cost of disproportionately more computational cost. For example, in order to increase the resolution of the grid by two, the number of grid cells increases by four in two dimensions, and by eight in three dimensions. For large or complicated domains, this results in a PDE that is very expensive to solve. The techniques we bring to bear are adaptive mesh refinement and parallelization. Adaptive mesh refinement allows for starting with a relatively coarse grid for simulation, then using successively finer grids in areas that need more precision. The idea of variable density meshes is widely used in research [1], and is used in many commercial products like Comsol [2] and Matlab [3]. This way the extra computation time is used only where it is needed. The use of parallel algorithms to solve PDEs has been proven useful in research [4][5], will further speed up the computation times of our tool with no cost to accuracy.

II. BACKGROUND

A. Finite Difference Method

The finite difference method is one of the simplest numerical PDE methods. This method aims to approximate derivatives over a short step, or difference, to get to an approximate solution. It divides the domain into a grid of discrete spatial points with a fixed distance between them. It then uses the distance between them as the difference needed to estimate derivatives.

The three most common variations are:

- Explicit Method
- Implicit Method
- Crank-Nicholson Method

B. Finite Element Method

The finite element method, similarly to the finite difference method, divides the domain into a collection of small elements that are easier to solve. This method, however, uses a triangular mesh rather than a square grid which makes it compatible with irregular spaces.

C. Multigrid Method

The multigrid method uses the solutions from coarse grids to augment the solution on a finer grid. This increases the convergence of a solution. This method can be paired with the finite element method as well.

D. Adaptive Grid

The rectangular grid of the finite difference method results in some unique challenges. Namely, curved or irregular boundaries can be difficult to discretize with a small error. If a boundary moves over time that can result in further error in due to the discretization errors. Traditionally the simplest solution to these problems is making the mesh much finer. Of course, a fine mesh leads to it's own complications. In higher dimensions we will see a n^2 or n^3 increase in computational time. It's also rather inefficient to use a fine mesh everywhere if most of the discretization errors occur in a single part of the grid.

Thus the proposed solution to this is a mesh grid that isn't the same resolution throughout. In areas where errors are occurring the grid resolution can be raised for more precision without having to do extra calculations in places where the lower resolution has acceptable error. This will better represent curves and changing boundaries will saving time and other computational resources.

III. METHODOLOGY

The finite difference method was chosen because the rectangular grid is simpler to break into pieces for parallel computation than the irregular grid of the finite element method. This will be implemented with an adaptive grid.

A. Project Overview

General flow of solution:

- 1) Get input of what the domain is. Basic rectangular, or a composition of rectangular shapes.
- 2) Get the boundary conditions at each section of the mesh.
- 3) Uniformly mesh the area, starting with some relative coarse mesh.

- 4) Set a maximum allowed discretization error

One way to do this is to find the maximum approximation error of the refined mesh, and any approx error on the coarse mesh higher than that value gets marked for refinement.

- 5) Evaluate approx discretization error at the nodes on the mesh. If they are above the threshold, then mark the node/edge for refinement.
- 6) Refine the marked edges. Can do this by expanding the grid. Multiple methods for this exist. I think for a rectangular mesh, we can just double the mesh precision.

- 7) Reevaluate on the refined grid to see if there are any points that still exceed the error threshold. re-expand.

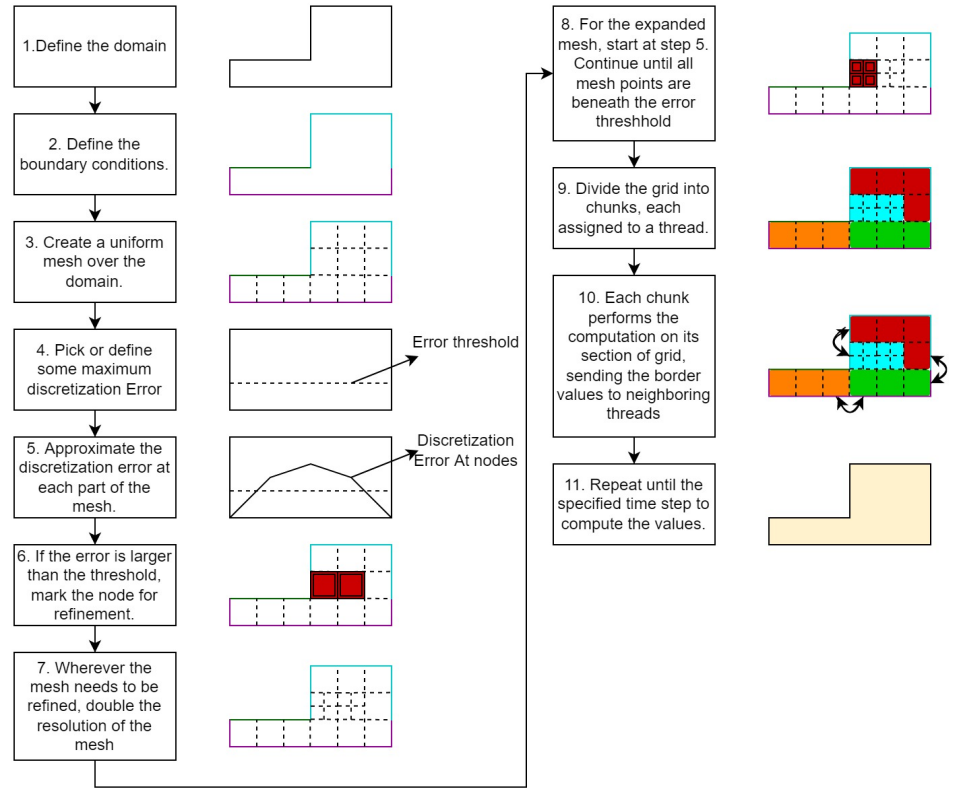


Fig. 1: Graphical overview of the general flow of the program

For the sake of efficiency this project operates off of a shared memory model. It uses a quadtree to determine how to adapt the grid and create a balanced amount of work for each process. The adaptive grid is stored as a sparse array.

B. The Quadtree




Fig. 2: The PR Quadtree [6]

To determine how the grid was going to be broken down a quadtree datastructure was used. This structure repeatedly divides space into four regions, and is an intuitive way to generate grids with variable resolutions.

Julia is missing some critical Quadtree functionality. To remedy that some custom functions were made.

1) *Nearest Neighbor*: This does exactly what it sounds like. It returns the nearest

neighbor and how far that neighbor is.

2) *Balance Quadtree*: Since the areas of finer mesh are likely to come from areas where the mesh has already been determined to not be sufficient it very easy for the quadtree to quickly become very unbalanced. The balance quadtree function fixes this.

C. Sparse Arrays

Fig. 3 shows a representation of a sparse array where the dots are data point and the blank spaces are areas where the data is either irrelevant or nonexistent. Storing and retrieving all of this irrelevant, or filler in the case of nonexistent, data is a waste of computational time and memory space. The fine areas of the mesh result in densely populated regions where the coarse areas are less populated. This makes a shared sparse array the ideal method for storing the variable mesh.

Julia is also missing some important functionality for using sparse arrays effectively. These functions had to be custom made to account for that.

1) *Closest Neighbor Value*: In a traditional array the closest neighbors are a simple array step away, but in sparse arrays the nearest actual value could be anywhere. Because of that there needs to be a way to both find what the value of the closest neighbor is, and how far away it is, so that it can then be pushed to the finite difference equation.

The function to handle that in this project not only returns the closest neighbor value and distance, but handles the instance where it turns out that there isn't a direct neighbor in a given direction. No neighbor indicated that a node is on the edge of a grid resolution change.

This is one of the most important custom functions, and it gets called every step of the parallel calculation process.



Fig. 3

IV. RESULTS AND DISCUSSION

The simulation was run with a constant number of 12 threads, but with a varying number of partitions. While there were 12 threads available to the program for each run have fewer partitions than threads is effectively using less threads and running things in a less parallel way.

Time to Execute, 12 Threads




Fig. 4: Results of the program run 12 threads and a varying number of partitions

As can be seen in Fig. 4, there was some degree of speedup when the program was run with more partitions. It isn't, however, to the degree that should be expected, not even two times speedup at any point. This is likely due to the expensive operations required to set up and then operate over the grid.

Also worth noting, Fig. 4 shows that program actually slows down when moving from 10 to 12 partition. At 12 partitions it is performing at the same speeds expected from 4 partitions which would be a drastic drop in performance if the gains had been greater. This too is likely caused by the operations required to set up the grid being expensive enough that the speedup could not make up for the increased complication of using 12 partitions.

As mentioned there is evidence to suggest that a fair amount of speedup without current methods could come from working out less costly ways of setting up and calculating each timestep.

In Fig. 5 it can be seen that other work using sparse arrays were able to iterate over many more elements in the same or less time than we were able to [7]. This implies that it should be more than possible to parallelize this problem using the methods lined out here with reasonably fast results. Once more, this points to the expensive custom functions as the likely culprit in the unsatisfactory results.

Grid Type	Max nodes	Runtime	Iterates	
			with < 5000 nodes	total
Predetermined	22148	6.00s	21	59
Boundary	15156	5.70s	25	67
Operator	15967	5.54s	33	71

Fig. 5: Results from "Adaptive Finite Difference Methods for Nonlinear Elliptic and Parabolic Partial Differential Equations with Free Boundaries" by Oberman and Zwiars

REFERENCES

- [1] Shinji Sakane, Tomohiro Takaki, and Takayuki Aoki. “Parallel-GPU-accelerated adaptive mesh refinement for three-dimensional phase-field simulation of dendritic growth during solidification of binary alloy”. In: *Materials Theory* 6.1 (Jan. 2022), p. 3. DOI: 10.1186/s41313-021-00033-5. URL: <https://doi.org/10.1186/s41313-021-00033-5>.
- [2] T Preney, J D Wheeler, and P Namy. “Adaptive mesh refinement: Quantitative computation of a rising bubble using COMSOL Multiphysics®”. en. In: (), p. 6.
- [3] *adaptmesh*. URL: <https://www.mathworks.com/help/pde/ug/adaptmesh.html>.
- [4] *Solve a 2D heat equation using data parallel C++*. URL: <https://www.intel.com/content/www/us/en/developer/articles/technical/solve-a-2d-heat-equation-using-data-parallel-c.html>.
- [5] Verena Horak and Peter Gruber. “Parallel Numerical Solution of 2-D Heat Equation”. en. In: (), p. 10.
- [6] *The PR Quadtree - CS3 Data Structures and Algorithms*. URL: <https://opensa-server.cs.vt.edu/ODSA/Books/CS3/html/PRquadtree.html>.
- [7] Adam M. Oberman and Ian Zwiars. “Adaptive Finite Difference Methods for Nonlinear Elliptic and Parabolic Partial Differential Equations with Free Boundaries”. en. In: *Journal of Scientific Computing* 68.1 (July 2016), pp. 231–251. DOI: 10.1007/s10915-015-0137-x. URL: <https://doi.org/10.1007/s10915-015-0137-x>.